

Yuchen Wang

✉ yuchenwang0303@gmail.com 📍 Mountain View, CA 🌐 yuchenwang3.github.io 🐙 yuchenwang3 🌐 yuchen3

Education

University of Illinois Urbana-Champaign , Computer Science	Champaign, IL, USA 2025 – 2027
Peking University , Intelligent Science and Technology (Artificial Intelligence) • Zhi Class (Elite Program), Top 10%	Beijing, China 2021 – 2025

Experience

Alibaba U.S. , Research Intern — LLM Training Algorithms & Infrastructure • Co-developed core RL post-training algorithms and engineered execution infrastructure for Occamy-1.0 , a 35B-A3B agent: captured exact model-boundary token IDs and verifier outcomes across OpenClaw, Hermes, and Accio Work. • Recovered long-horizon SFT/RL supervision through workspace replay and context-action-reward realignment after tool failures and compaction, keeping trajectories usable for post-training. • Shipped Occamy with Accio; versus its Qwen3.6 base, the released model raised AutomationBench strict pass from 7.5% to 27.6% and Claw-Eval average from 69.5 to 82.2, while reducing Claw-Eval T/C full-trial time by 36.6%.	Sunnyvale, CA, USA May 2026 – Dec 2026 8 months
Freedo Technology , Research Intern — 3D Vision and Generative AI • Built an end-to-end depth-conditioned ControlNet pipeline for 3D reconstruction from noisy point clouds, improving geometric fidelity from 32.7% to 85.3%.	Beijing, China May 2025 – Aug 2025 4 months

Open Source Projects

Open-Source LLM Systems Engineering

Upstream reliability and performance work across Hugging Face Datasets, mcore-bridge, FlashInfer, vLLM, NVIDIA NeMo, Megatron-LM, and agent runtimes.

- Data and packed attention: preserved source sharding in [HF Datasets streaming shuffle](#), restoring four-worker use on 19 shards. Made packed QSA selection document-local in [mcore-bridge](#), 2.70× faster in an 8K synthetic selector measurement.
- Numerics and memory: in-place BF16 MoE conversion in [vLLM](#) halved TP2 conversion peak allocation from 7.88 to 3.94 GiB; fixed calibrated FP8 KV prefill in [FlashInfer](#). Fused GDN Q/K normalization and added selective Mamba recompute in Megatron-LM.
- Distributed training: deferred NeMo RL teacher top-k payloads for a 4.44–5.31× speedup in controlled Ray transfer measurements. Developing [Bridge-local Qwen4-Exp](#) support for QSA, per-layer embeddings, gated residuals, and HF conversion; GPU integration remains pending.
- Reliability: fixed small-norm Muon normalization, proposed checkpoint-safe Adam counters for Megatron-LM, and co-authored a NeMo Gym fix preserving HTTP errors across process boundaries. Repaired tool-schema coercion and SSH reconnect in Hermes.

Projects

RareDx — Graph-Grounded RL for Rare-Disease Diagnosis

Preprint under review; advised by Prof. Tianyu Liu at Tsinghua University. Graph-grounded rewards for partial diagnostic matches.

- Co-designed rewards using disease-ontology distance, biomedical similarity, and phenotype-disease links. With SFT and retrieval, 27B diagnosis Hit@10 increased from 28.53% to 40.76%.

CineFlow — Dependency-Driven Parallel Video Generation

Dependency-driven video diffusion using semantic DAGs, runtime trajectory fusion, online edge pruning, and critical-path scheduling without retraining. Advised by Prof. Fan Lai at UIUC.

- Analyzed 1K+ MovieGenBench prompts; more than 68% exhibit non-linear dependencies.
- Across Wan2.2-5B, CogVideoX-5B, and HunyuanVideo on 8×H100, achieved 1.7–5.5× end-to-end speedup, 1.30–2.02× lower P90 latency, and 5.4–17.3% higher VBench overall.

Dynamic Prefill Optimization via Adaptive Online Packing

AIMD control with live p95 TTFT feedback and greedy/DP prompt packing.

- Achieved up to 20% lower TTFT on production-style DynamoLLM traces.

FlashAttention-style CUDA Optimization for GPT-2

Tiled online softmax, memory-efficient attention, and kernel fusion.

- Reduced HBM traffic by roughly 10× and achieved up to 9% end-to-end speedup.

Reinforcement Learning for Legal Reasoning

Advised by Prof. Yansong Feng at Peking University; Zero-RL → distilled-CoT SFT → GRPO for structured legal multiple-choice reasoning.

- Reached 57.6% accuracy, outperforming SFT-only and RL-only baselines.

Academic Service

- Reviewer, [WSDM 2027](#).

Awards

Merit Student

Zhi Class Scholarship

CMO (Provincial First Prize), 2020

Skills

- Python, C++, CUDA, PyTorch, Ray, FSDP, Megatron-LM, NeMo RL, vLLM, SGLang, Nsight, RL, distributed training, diffusion transformers, ControlNet, 3D vision, and point clouds