

Yuchen Wang

✉ yuchenwang0303@gmail.com 📍 Mountain View, CA 🌐 yuchenwang3.github.io 🐙 yuchenwang3 🌐 yuchen3

Education

University of Illinois Urbana-Champaign , Computer Science	Champaign, IL, USA 2025 – 2027
Peking University , Intelligent Science and Technology (Artificial Intelligence) • Zhi Class (Elite Program), Top 10%	Beijing, China 2021 – 2025

Experience

Alibaba U.S. , Research Intern — Agentic LLM Post-Training	Sunnyvale, CA, USA May 2026 – present 5 months
- Core contributor to Occamy-1.0 (arXiv · Website), a 35B-A3B agent model. Built execution-grounded data and post-training infrastructure for long-horizon tool use; the team's SFT corpus contains 15K trajectories and 403M tokens. - Built verifier-gated pipelines for OpenClaw, Hermes, and Accio Work: token-exact replay, state recovery, and compaction-aware segments preserving tokens, log-probs, and loss masks for SFT/RL. - The released model improves Claw-Eval average from 69.5 to 82.2 and AutomationBench strict pass from 7.5% to 27.6%, reaching the Pareto frontier in the report's four-benchmark cost comparison.	
Freedo Technology , Research Intern — 3D Vision and Generative AI	Beijing, China May 2025 – Aug 2025 4 months
Built an end-to-end depth-conditioned ControlNet pipeline for 3D reconstruction from noisy point clouds, improving geometric fidelity from 32.7% to 85.3%.	

Open Source Projects

Open-Source LLM Systems Engineering

Upstream reliability and performance work across FlashInfer, vLLM, NVIDIA NeMo, Megatron-LM, Hermes Agent, SGLang, verl, vLLM/Vime, and ModelScope.

- Numerics and training: fixed FP8 KV calibration in FlashInfer and small-norm Muon normalization in Emerging Optimizers; fused GatedDeltaNet Q/K normalization and added selective Mamba recompute in Megatron-LM.
- Performance and scale: in-place BF16 MoE conversion in vLLM halved TP2 conversion peak allocation from 7.88 to 3.94 GiB. Deferred teacher top-k payloads in NeMo RL achieved 4.44–5.31× speedup in controlled Ray transfer measurements.
- Agent runtime: repaired schema-reference argument coercion and SSH teardown/reconnect in Hermes; co-authored pickle-safe cross-process HTTP failure propagation in NeMo Gym.

Projects

CineFlow — Dependency-Driven Parallel Video Generation

Dependency-driven video diffusion using semantic DAGs, runtime trajectory fusion, online edge pruning, and critical-path scheduling without retraining. Advised by Prof. Fan Lai at UIUC.

- Analyzed 1K+ MovieGenBench prompts; more than 68% exhibit non-linear dependencies.
- Across Wan2.2-5B, CogVideoX-5B, and HunyuanVideo on 8×H100, achieved 1.7–5.5× end-to-end speedup, 1.30–2.02× lower P90 latency, and 5.4–17.3% higher VBench overall.

Dynamic Prefill Optimization via Adaptive Online Packing

AIMD control with live p95 TTFT feedback and greedy/DP prompt packing.

- Achieved up to 20% lower TTFT on production-style DynamoLLM traces.

FlashAttention-style CUDA Optimization for GPT-2

Tiled online softmax, memory-efficient attention, and kernel fusion.

- Reduced HBM traffic by roughly 10× and achieved up to 9% end-to-end speedup.

Reinforcement Learning for Legal Reasoning

Advised by Prof. Yansong Feng at Peking University; Zero-RL → distilled-CoT SFT → GRPO for structured legal multiple-choice reasoning.

- Reached 57.6% accuracy, outperforming SFT-only and RL-only baselines.

Skills

- Python, C++, CUDA, PyTorch, Ray, FSDP, Megatron-LM, NeMo RL, vLLM, SGLang, Nsight, RL, distributed training, diffusion transformers, ControlNet, 3D vision, and point clouds