Yuchen Wang

Agentic post-training | ML systems

yuchenwang0303@gmail.com | (217) 480-2166 | Website | GitHub | LinkedIn | WeChat

Education

University of Illinois Urbana-Champaign 2025 – 2027 (expected)
Master of Computer Science Champaign, IL

Peking University 2021 – 2025
B.S. in Intelligent Science and Technology (AI), Zhi Class. Top 10%. Beijing, China

Experience

Alibaba U.S. – Accio Team May 2026 – present
Research Intern, LLM Training Algorithms & Infrastructure Sunnyvale, CA

- Core contributor to **Occamy-1.0** [report]. Compared with its Qwen3.6-35B-A3B base, the team's model raised AutomationBench strict pass rate from **7.5% to 27.6%** and Claw-Eval average score from **69.5 to 82.2**. Full trials on the Claw-Eval T/C subset took 36.6% less time.
- Co-designed staged SFT/HDPO, model soup, and SAO for long-horizon tool use, with trainable context compaction and cross-segment credit assignment.
- Designed task generators and graders that check the environment after an agent acts. Built asynchronous RL infrastructure for OpenClaw, Hermes, and Accio Work. The team's SFT dataset spans **15K trajectories and 403M tokens**.

🤗 800K+ downloads across official Occamy releases in the 30 days to Nov. 1, 2026.

Freedo Technology May – Aug 2025
Research Intern, 3D Vision and Generative AI Beijing, China

- Improved point-cloud reconstruction fidelity from **32.7% to 85.3%** by conditioning ControlNet on depth and surface normals to preserve edges and planar surfaces.

Selected Open-Source Contributions

- **Hybrid training (Megatron-LM).** Added selective Mamba recomputation and in-kernel GDN Q/K normalization to reduce stored activations. Mamba forward/backward parity tests passed on H100 with TP1, TP2, and sequence parallelism.
- **RL distillation (NeMo RL).** Bypassed Ray driver materialization with DP-local teacher top-k handoff, achieving **4.44–5.31× faster transfers** in controlled 8-B200 prototype benchmarks.
- **Optimizer correctness.** Restored Muon scale invariance for small gradients, isolated GDN gradient clipping, and preserved Adam step counters across distributed checkpoint restores.
- **Model integration (Megatron Bridge).** Contributed Qwen4-Exp text-decoder support spanning document-local QSA, per-layer embedding conversion, bounded RMSNorm backward, and low-precision gate rounding.
- **Inference memory (vLLM).** Converted BF16 MoE weights in place, cutting peak allocation from **7.88 to 3.94 GiB** in an exact-shape TP2 benchmark for a 120B model.

Selected Projects

University of Illinois Urbana-Champaign, 2025 – 2026

- Adaptive prompt packing. Cut time to first token by **up to 20%** on production-style DynamoLLM traces by adjusting batch size from live latency feedback and packing prompts by length.
- GPT-2 CUDA kernels. Tiled attention and fused kernels cut HBM traffic roughly **10×** and raised arithmetic intensity from 15 to 157 FLOP/byte. End-to-end speed improved by up to 9%.

Technical Skills

Reviewer, WSDM 2027

Python, C++, CUDA, PyTorch, Ray, FSDP, Megatron-LM, NeMo RL, vLLM, SGLang, Nsight.