

Yuchen Wang

(217) 480-2166 | yuchenwang0303@gmail.com | [LinkedIn](#) | [GitHub](#) | [Portfolio](#)

EDUCATION

University of Illinois Urbana-Champaign (UIUC)

M.S. in Computer Science

2025 – 2027 (expected)

Champaign, IL, USA

Peking University

B.S. in Intelligent Science and Technology (Artificial Intelligence) Zhi Class (Elite Program), Top 10%

2021 – 2025

Beijing, China

INTERNSHIP EXPERIENCE

• Alibaba U.S.

Research Intern — Agentic LLM Post-Training

May 2026 – Dec 2026

Sunnyvale, CA, USA

- Contributed to **Occamy-1.0** with Alibaba's Accio Team [[Code](#)], focusing on data and post-training infrastructure. The Qwen3.6-35B-A3B-based recipe combines Marathon (SFT + HDPO) and Sprint (SFT) experts, model soup, and post-merge SAO reinforcement learning.
- Built verifier-gated multi-harness data pipelines with token-exact proxy/TiTO replay, state reconstruction, provenance, context-segment lineage, and loss masks. On matched Claw-Eval T/C, the model uses **19.5% fewer tokens/trajectory** and **46.4% less trace wall time** than the starting checkpoint; timeout rate falls from 9.88% to 2.18%.
- Latest report: **82.2 Claw-Eval average (71.4 Pass³)**, **49.16 WildClawBench**, and **27.6% AutomationBench strict pass rate** (team-level model results).

• Freedo Technology

Research Intern — 3D Vision and Generative AI

May 2025 – Aug 2025

Beijing, China

- Built an end-to-end depth-conditioned ControlNet pipeline for 3D building reconstruction from noisy point clouds; designed depth/normal conditioning to preserve structural edges and planar surfaces, improving geometric fidelity from **32.7% to 85.3%**.

OPEN-SOURCE SYSTEMS ENGINEERING

- Reliability & correctness:** Muon, hybrid Mamba/MoE, packing, NCCL-startup, and optimizer fixes across [NVIDIA NeMo](#), [vLLM/Vime](#), and four [ms-swift](#) patches; co-authored [NeMo Gym #2726](#) for pickle-safe cross-process HTTP failures. Current work covers non-finite RL logprobs ([NeMo RL #2962](#)), GDN/Muon clipping and routing ([Megatron-LM #5395/#5400/#5431](#)), skipped-tensor weight checks ([SGLang #31621](#)), and authoritative FSDP strategy ([verl #7597](#)).
- Performance & scale:** [vLLM #54699](#) converts FlashInfer BF16 MoE weights in-place, cutting TP2 peak allocation from **7.88 to 3.94 GiB** and enabling real 120B load/generation plus 3/3 sleep-wake cycles on 8×B200; [NeMo RL #3943](#) delivers **4.44–5.31× controlled Ray transfer speedup** while reducing driver RSS deltas to 0.1–0.3 MB; [Megatron-LM #5396/#5463](#) fuse GatedDeltaNet normalization and add Mamba recompute for long-context training.

PAPERS & PROJECTS

• CineFlow: Dependency-Driven Parallel Video Generation [[Paper](#)]

Oct 2025 – May 2026

Advisor: Prof. Fan Lai, University of Illinois Urbana-Champaign

Tools: PyTorch, xDiT, Diffusion Transformers

- Architected CineFlow, the first dependency-driven video-diffusion inference system, after analyzing 1K+ MovieGenBench prompts (over 68% exhibit non-linear dependencies); compiled prompts into semantic DAGs, fused runtime trajectories with online edge pruning, and scheduled critical paths without retraining.
- Evaluated Wan2.2-5B, CogVideoX-5B, and HunyuanVideo on 8×H100, achieving **1.7–5.5× end-to-end speedup**, **1.30–2.02× lower P90 latency**, and **5.4–17.3% higher VBench overall**; pruning added nearly 7% throughput and improved visual coherence by 6.7–14.6%.

• Dynamic Prefill Optimization for LLM Inference [[Report](#)]

Aug 2025 – Jan 2026

Adaptive Online Packing — University of Illinois Urbana-Champaign

Tools: Python, CUDA, vLLM, Profiling

- Built an AIMD batching controller with real-time p95 TTFT feedback and greedy/DP length-aware packing before continuous batching; achieved up to **20% lower TTFT** on production-style DynamoLLM traces.

• FlashAttention-style CUDA Kernels for GPT-2 [[Report](#)]

Aug 2025 – Dec 2025

University of Illinois Urbana-Champaign

Tools: CUDA, C++, Nsight Compute/Systems

- Implemented tiled attention with online softmax and kernel fusion, raising arithmetic intensity from 15 to 157 FLOP/byte and delivering **~10× lower HBM traffic** and up to **~9% end-to-end speedup**.

• RL for Legal Reasoning on Multi-Choice QA [[Thesis](#)]

Jan 2025 – Mar 2025

Advisor: Prof. Yansong Feng, Peking University

Tools: PyTorch, Qwen, GRPO

- Built a Zero-RL → distilled-CoT SFT → GRPO pipeline for structured legal reasoning, reaching **57.6% accuracy** and outperforming SFT-only and RL-only baselines.

SKILLS

Python, C++, CUDA, PyTorch, Ray, FSDP, 3D vision, point clouds, ControlNet, diffusion transformers, Megatron-LM, NeMo-RL, vLLM, SGLang, distributed training, profiling